

Exercise 8.2 - Time Series Modeling

using the dataset `us_retail_sales.csv` for this assignment. This data gives the total monthly retail sales in the US from January 1992 until June 2021. With this dataset, complete the following steps:

1. Plot the data with proper labeling and make some observations on the graph.
2. Split this data into a training and test set. Use the last year of data (July 2020 – June 2021) of data as your test set and the rest as your training set.
3. Use the training set to build a predictive model for the monthly retail sales.
4. Use the model to predict the monthly retail sales on the last year of data.
5. Report the RMSE of the model predictions on the test set.

```
In [1]: ## import the required packages  
  
import pandas as pd  
import numpy as nm  
import matplotlib.pyplot as plt  
from sklearn.linear_model import LinearRegression  
from sklearn import metrics  
from datetime import datetime
```

```
In [2]: ## import the retail sales data into dataframe  
  
sales_df = pd.read_csv('us_retail_sales.csv')  
sales_df.head()
```

```
Out[2]:
```

| | YEAR | JAN | FEB | MAR | APR | MAY | JUN | JUL | AUG | SEP | OCT | |
|---|------|--------|--------|--------|--------|--------|--------|----------|----------|----------|----------|-----|
| 0 | 1992 | 146925 | 147223 | 146805 | 148032 | 149010 | 149800 | 150761.0 | 151067.0 | 152588.0 | 153521.0 | 153 |
| 1 | 1993 | 157555 | 156266 | 154752 | 158979 | 160605 | 160127 | 162816.0 | 162506.0 | 163258.0 | 164685.0 | 166 |
| 2 | 1994 | 167518 | 169649 | 172766 | 173106 | 172329 | 174241 | 174781.0 | 177295.0 | 178787.0 | 180561.0 | 180 |
| 3 | 1995 | 182413 | 179488 | 181013 | 181686 | 183536 | 186081 | 185431.0 | 186806.0 | 187366.0 | 186565.0 | 189 |
| 4 | 1996 | 189135 | 192266 | 194029 | 194744 | 196205 | 196136 | 196187.0 | 196218.0 | 198859.0 | 200509.0 | 200 |

Understanding the data

```
In [3]: sales_df.shape
```

```
Out[3]: (30, 13)
```

```
In [4]: sales_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 30 entries, 0 to 29
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   YEAR        30 non-null    int64
1   JAN         30 non-null    int64
2   FEB         30 non-null    int64
3   MAR         30 non-null    int64
4   APR         30 non-null    int64
5   MAY         30 non-null    int64
6   JUN         30 non-null    int64
7   JUL         29 non-null    float64
8   AUG         29 non-null    float64
9   SEP         29 non-null    float64
10  OCT         29 non-null    float64
11  NOV         29 non-null    float64
12  DEC         29 non-null    float64
dtypes: float64(6), int64(7)
memory usage: 3.2 KB

```

Creating a one dimensional data set of the sales data set with Year being the ID, Month being the variable and the sales being the value

```

In [5]: sales_data = pd.melt(sales_df, id_vars=["YEAR"], var_name="Month", value_name="Sales")
sales_data

```

```

Out[5]:
   YEAR  Month  Sales
0  1992   JAN  146925.0
1  1993   JAN  157555.0
2  1994   JAN  167518.0
3  1995   JAN  182413.0
4  1996   JAN  189135.0
...
355  2017  DEC  433282.0
356  2018  DEC  434803.0
357  2019  DEC  458055.0
358  2020  DEC  484782.0
359  2021  DEC      NaN

```

360 rows × 3 columns

Replacing the month names with numbers

```

In [6]: sales_data["Month"].replace({'JAN': 1, 'FEB': 2, 'MAR': 3, 'APR': 4, 'MAY': 5, 'JUN': 6,
    'SEP': 9, 'OCT': 10, 'NOV': 11, 'DEC': 12}, inplace=True)
sales_data

```

```

Out[6]:
   YEAR  Month  Sales

```

| | YEAR | Month | Sales |
|-----|------|-------|----------|
| 0 | 1992 | 1 | 146925.0 |
| 1 | 1993 | 1 | 157555.0 |
| 2 | 1994 | 1 | 167518.0 |
| 3 | 1995 | 1 | 182413.0 |
| 4 | 1996 | 1 | 189135.0 |
| ... | ... | ... | ... |
| 355 | 2017 | 12 | 433282.0 |
| 356 | 2018 | 12 | 434803.0 |
| 357 | 2019 | 12 | 458055.0 |
| 358 | 2020 | 12 | 484782.0 |
| 359 | 2021 | 12 | NaN |

360 rows × 3 columns

creating a new field Date based on the Year and Month fields and assigning the first day of that month.

```
In [7]: sales_data['DATE'] = pd.to_datetime(sales_data[['YEAR', 'Month']].assign(DAY=1))
sales_data
```

```
Out[7]:
```

| | YEAR | Month | Sales | DATE |
|-----|------|-------|----------|------------|
| 0 | 1992 | 1 | 146925.0 | 1992-01-01 |
| 1 | 1993 | 1 | 157555.0 | 1993-01-01 |
| 2 | 1994 | 1 | 167518.0 | 1994-01-01 |
| 3 | 1995 | 1 | 182413.0 | 1995-01-01 |
| 4 | 1996 | 1 | 189135.0 | 1996-01-01 |
| ... | ... | ... | ... | ... |
| 355 | 2017 | 12 | 433282.0 | 2017-12-01 |
| 356 | 2018 | 12 | 434803.0 | 2018-12-01 |
| 357 | 2019 | 12 | 458055.0 | 2019-12-01 |
| 358 | 2020 | 12 | 484782.0 | 2020-12-01 |
| 359 | 2021 | 12 | NaN | 2021-12-01 |

360 rows × 4 columns

Sorting the dataset on date and resetting the index values

```
In [8]: sales_data = sales_data.sort_values(by=['DATE'])
```

```
sales_data
```

Out[8]:

| | YEAR | Month | Sales | DATE |
|-----|------|-------|----------|------------|
| 0 | 1992 | 1 | 146925.0 | 1992-01-01 |
| 30 | 1992 | 2 | 147223.0 | 1992-02-01 |
| 60 | 1992 | 3 | 146805.0 | 1992-03-01 |
| 90 | 1992 | 4 | 148032.0 | 1992-04-01 |
| 120 | 1992 | 5 | 149010.0 | 1992-05-01 |
| ... | ... | ... | ... | ... |
| 239 | 2021 | 8 | NaN | 2021-08-01 |
| 269 | 2021 | 9 | NaN | 2021-09-01 |
| 299 | 2021 | 10 | NaN | 2021-10-01 |
| 329 | 2021 | 11 | NaN | 2021-11-01 |
| 359 | 2021 | 12 | NaN | 2021-12-01 |

360 rows × 4 columns

In [9]:

```
sales_data = sales_data.reset_index(drop=True)
sales_data
```

Out[9]:

| | YEAR | Month | Sales | DATE |
|-----|------|-------|----------|------------|
| 0 | 1992 | 1 | 146925.0 | 1992-01-01 |
| 1 | 1992 | 2 | 147223.0 | 1992-02-01 |
| 2 | 1992 | 3 | 146805.0 | 1992-03-01 |
| 3 | 1992 | 4 | 148032.0 | 1992-04-01 |
| 4 | 1992 | 5 | 149010.0 | 1992-05-01 |
| ... | ... | ... | ... | ... |
| 355 | 2021 | 8 | NaN | 2021-08-01 |
| 356 | 2021 | 9 | NaN | 2021-09-01 |
| 357 | 2021 | 10 | NaN | 2021-10-01 |
| 358 | 2021 | 11 | NaN | 2021-11-01 |
| 359 | 2021 | 12 | NaN | 2021-12-01 |

360 rows × 4 columns


In [10]:

```
## Plotting the dataset - date vs retail sales

x = sales_data.DATE
y = sales_data.Sales
```

```
plt.plot(x, y)
plt.xlabel("Date")
plt.ylabel("Retail Sales")
plt.title("Sales per Year")
```

Out[10]: Text(0.5, 1.0, 'Sales per Year')



We can see there was a slump in the retail sales during the 2008-2009 recession, but it slowly picked up until the 2020 pandemic early that year which again saw a steep fall for a brief period of time.

```
In [11]: ## Dropping all the nulls in the dataset

sales_data.dropna(inplace=True)
sales_data
```

Out[11]:

| | YEAR | Month | Sales | DATE |
|-----|------|-------|----------|------------|
| 0 | 1992 | 1 | 146925.0 | 1992-01-01 |
| 1 | 1992 | 2 | 147223.0 | 1992-02-01 |
| 2 | 1992 | 3 | 146805.0 | 1992-03-01 |
| 3 | 1992 | 4 | 148032.0 | 1992-04-01 |
| 4 | 1992 | 5 | 149010.0 | 1992-05-01 |
| ... | ... | ... | ... | ... |
| 349 | 2021 | 2 | 504458.0 | 2021-02-01 |
| 350 | 2021 | 3 | 559871.0 | 2021-03-01 |
| 351 | 2021 | 4 | 562269.0 | 2021-04-01 |
| 352 | 2021 | 5 | 548987.0 | 2021-05-01 |
| 353 | 2021 | 6 | 550782.0 | 2021-06-01 |

354 rows × 4 columns

In [12]:

```
## adding a new feature OrdDate which is the integer conversion of the date value for t

sales_data['OrdDate'] = pd.to_datetime(sales_data['DATE'])
sales_data['OrdDate'] = sales_data['OrdDate'].map(datetime.toordinal)
```

Split this data into a training and test set. Use the last year of data (July 2020 – June 2021) of data as your test set and the rest as your training set.

Sales is the target value to be predicted.

```
In [13]: ## splitting data into train and test sets

train_df = sales_data[sales_data['DATE'] < '2020-07-01']
test_df = sales_data[sales_data['DATE'] >= '2020-07-01']
```

```
In [14]: X_train = train_df[['OrdDate']]
y_train = train_df[['Sales']]
X_test = test_df[['OrdDate']]
y_test = test_df[['Sales']]
```

```
In [15]: X_test.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 12 entries, 342 to 353
Data columns (total 1 columns):
 #   Column   Non-Null Count  Dtype  
---  -
 0   OrdDate  12 non-null     int64  
dtypes: int64(1)
memory usage: 192.0 bytes
```

Use the training set to build a predictive model for the monthly retail sales.

Building a prediction model using Linear Regression model

```
In [16]: ## Linear regression modeling
# Create a model
model = LinearRegression()

# Fit the model to the training set
model.fit(X_train, y_train)
```

```
Out[16]: LinearRegression()
```

```
In [17]: ## predicting the sales values on the test set

test_predictions = model.predict(X_test)
test_predictions
```

```
Out[17]: array([[449450.188174  ],
```

```
[450339.37662852],  
[451228.56508304],  
[452089.07003903],  
[452978.25849355],  
[453838.76344953],  
[454727.95190405],  
[455617.14035857],  
[456420.2783175 ],  
[457309.46677202],  
[458169.971728 ],  
[459059.16018252]])
```

```
In [18]: y_test = y_test.reset_index(drop=True)  
y_test
```

```
Out[18]:
```

| | Sales |
|----|----------|
| 0 | 481627.0 |
| 1 | 483716.0 |
| 2 | 493327.0 |
| 3 | 493991.0 |
| 4 | 488652.0 |
| 5 | 484782.0 |
| 6 | 520162.0 |
| 7 | 504458.0 |
| 8 | 559871.0 |
| 9 | 562269.0 |
| 10 | 548987.0 |
| 11 | 550782.0 |

```
In [19]: ## converting the prediction result array to dataframe  
  
pred_test = pd.DataFrame(test_predictions, columns=['pred_sales'])  
pred_test
```

```
Out[19]:
```


| | pred_sales |
|---|---------------|
| 0 | 449450.188174 |
| 1 | 450339.376629 |
| 2 | 451228.565083 |
| 3 | 452089.070039 |
| 4 | 452978.258494 |
| 5 | 453838.763450 |
| 6 | 454727.951904 |
| 7 | 455617.140359 |

| | pred_sales |
|----|---------------|
| 8 | 456420.278317 |
| 9 | 457309.466772 |
| 10 | 458169.971728 |
| 11 | 459059.160183 |

```
In [20]: ## plotting the actuals and predicted values.

plt.plot(y_test, label = 'Actuals')
plt.plot(test_predictions, label = 'Predictions')
plt.legend()
```

Out[20]: <matplotlib.legend.Legend at 0x1fd511bcf10>



```
In [21]: print('Test RMSE:', metrics.mean_squared_error(y_test, test_predictions, squared=False))
```

Test RMSE: 66429.10224838056

From the Actuals and Predictions plot and RMSE score, we can see the predicted values are way off the actuals. This could be because the actuals had a spike during the mid of 2020, during the pandemic.